


Milestone 5: Model Evaluation & Analysis

Project: Talk to Your Database - A Natural-Language Analytics Copilot



Course: Data Science and AI Lab

Submission date: August 6, 2026

Team Members (Group 10):

1. Siddhant Hitesh Mantri (21f3002218)
2. Anirudh Komanduri (22f1000522)
3. Vishal S (23f2003089)
4. Walunila Aier (21f3002564)
5. Sambhav Jha (22f3003227)

Table of Contents

Executive Summary.....	4
1. Introduction.....	4
1.1 Project Objective.....	4
1.2 Scope of Evaluation.....	4
1.3 Models Evaluated.....	4
1.4 Objectives of Milestone 5.....	4
2. Experimental Setup.....	5
2.1 Hardware and Software Environment.....	5
2.2 Frameworks and Libraries Used.....	5
2.3 Dataset Versions.....	5
2.4 Training/Validation/Test Split.....	5
2.5 Random Seeds.....	5
2.6 Reproducibility Details.....	6
2.7 Evaluation Workflow.....	6
3. Model Training Summary.....	6
3.1 Final Model Architecture.....	6
3.2 Hyperparameters.....	6
3.3 Checkpoint Selection.....	7
3.4 Input Representation.....	7
4. Evaluation Methodology.....	7
4.1 Evaluation Protocol.....	7
4.2 Test Dataset Description.....	7
4.3 Ground Truth Preparation.....	7
4.4 Baseline Used.....	7
4.5 Success Criteria.....	7
5. Performance Metrics.....	7
6. Experimental Results.....	8
6.1 Prompt Configuration Ablation.....	8
6.2 Prediction Execution Status (M-Schema, Primary).....	8
7. Baseline Comparison.....	8
8. Hyperparameter Analysis.....	8
8.1 M-Schema vs. DDL - Training-Time Finding (Spider, Milestone 3).....	8
8.2 Independent Cross-Benchmark Evidence (This Milestone).....	9
9. Ablation Study.....	9
9.1 Prompt Format: DDL vs. M-Schema.....	9
9.2 Available Context: M-Schema vs. M-Schema + Evidence.....	9
10. Error Analysis.....	10
10.1 Error Taxonomy (M-Schema, Primary) - What Failed, and Why.....	10
10.2 Feature-Conditioned Accuracy - What Makes a Query Harder.....	11
10.3 Testing the Predicted Mechanisms.....	11

10.4 Schema-Linking Precision.....	12
10.5 The 3-JOIN Anomaly: A Database-Level Effect, Not a JOIN-Count Effect.....	12
10.6 All Three Configurations Compared.....	13
Representative examples (M-Schema).....	13
10.7 Post-Hoc Repair - A Negative Result.....	14
10.8 Concrete, Evidence-Based Improvements.....	14
11. Model Robustness.....	14
11.1 Generalization and Edge Cases.....	14
11.2 Invalid, Ambiguous, and Out-of-Domain Requests.....	14
12. Computational Performance.....	15
13. Limitations.....	15
14. Possible Improvements.....	15
15. Discussion.....	15
15.1 Were the Objectives Achieved?.....	15
15.2 Tying Milestone 2's EDA to Milestone 5's Errors.....	16
15.3 Milestone 3 and 4 Decisions, Revisited.....	16
15.4 Assumptions Validated or Disproved.....	16
16. Conclusion.....	16
16.1 Summary of Findings.....	16
16.2 Deployment Readiness - Quantified.....	17
16.3 Readiness for Further Work.....	17
Appendix.....	18
A. Evaluation Scripts.....	18
B. Full Metric Files.....	18
C. Dataset.....	18
D. Colab Notebook for Milestone 5.....	18
Signature of Approval.....	19

Executive Summary

Milestone 4 selected and trained checkpoint 375, a QLoRA adapter for Qwen/Qwen3-4B-Instruct-2507, achieving 85.514% strict execution accuracy on internal tuning data, and every training example used the M-Schema prompt format.

This report evaluates checkpoint 375 on BIRD Mini-Dev (500 examples), a benchmark never used in training or model selection. Spider validation is not used for evaluation here: Milestone 4 itself identifies that split as unsuitable for an unbiased final claim, since it was examined during Milestone 3 selection, and calls for BIRD as the appropriate benchmark for this milestone.

Checkpoint 375 was evaluated under three prompt configurations on the identical 500 questions: plain DDL schema (16.8% strict execution accuracy), M-Schema - the format used in training (25.8%, primary result), and M-Schema with BIRD's evidence hints added (30.8%, the highest verified accuracy for this checkpoint). M-Schema without evidence is designated the primary result because it most closely matches training conditions: Spider, the training data, has no evidence field, so this configuration isolates generalization ability without introducing help the model never had during training. These two configuration changes are independent and additive (+9.0 and +5.0 points respectively), showing that evaluation configuration accounts for a meaningful share of the apparent generalization gap, though a substantial gap remains under any configuration.

Beyond accuracy, this report contributes: a causal error taxonomy testing (not assuming) the mechanism behind each failure category (Section 10.1–10.3); a schema-linking precision analysis showing the model references real tables 100% of the time but real columns only 82.8% of the time (Section 10.4); an 18-prompt adversarial robustness test showing the system stays read-only under explicit write and prompt-injection attempts (Section 11.2); and a quantified, context-specific deployment-readiness assessment (Section 16.2).

1. Introduction

1.1 Project Objective

The project converts a natural-language question and a database schema into a single, safe, read-only SQL query, executes it through an immutable connection, and returns both the SQL and result - the SQL remains visible so a human can inspect it before trusting the output.

1.2 Scope of Evaluation

Milestone 5 evaluates the frozen adapter from Milestone 4 (checkpoint 375, label milestone4_frozen) on BIRD Mini-Dev, held out from all training and selection decisions. Spider validation is not used: Milestone 4 identifies it as unsuitable for an unbiased final claim on this checkpoint, and BIRD is its own recommended substitute. No further training or checkpoint selection occurs in this milestone.

1.3 Models Evaluated

Qwen/Qwen3-4B-Instruct-2507 (4.022B parameters) + checkpoint 375 QLoRA adapter (33.03M trainable parameters, 0.821% of base).

1.4 Objectives of Milestone 5

Establish a defensible, held-out generalization estimate using BIRD Mini-Dev.

Produce a causal error taxonomy, testing rather than assuming each failure mechanism.

Quantify the effect of prompt format and available context as controlled ablations.

Synthesize findings across Milestones 2–4, not treat evaluation as isolated.

Provide a quantified, context-specific deployment-readiness assessment.

2. Experimental Setup

2.1 Hardware and Software Environment

Component	Value
GPU	NVIDIA L4 (22.03 GiB)
torch / transformers	2.11.0+cu128 / 5.13.1
torchao (pinned)	0.18.0 - 0.10.0 is incompatible with the peft version used for adapter loading
bitsandbytes (pinned)	0.48.1
Batch size	8 (evaluation); resume-safe against interruption

2.2 Frameworks and Libraries Used

- PyTorch (2.11.0+cu128) - base tensor/training framework
- Transformers (5.13.1) - model loading, tokenization, generation
- PEFT - QLoRA adapter loading (checkpoint 375)
- bitsandbytes (0.48.1, pinned) - 4-bit quantized inference
- torchao (0.18.0, pinned) - required for compatibility with the peft version used for adapter loading; 0.10.0 is incompatible
- sqlglot - SQL parsing, syntax validation, AST-based table/column reference extraction (error taxonomy, schema-linking precision)
- SQLite3 (Python stdlib) - read-only, timeout-protected query execution against BIRD databases
- pandas / matplotlib / seaborn - analysis and visualization

2.3 Dataset Versions

- BIRD Mini-Dev (500 examples, 11 SQLite databases) - held out from all training and selection; used for evaluation. (*Appendix C*)
- Spider 1.0 validation (1,034 examples) - not used for evaluation (Section 1.2); mentioned only where Milestone 3/4's own figures are cited for context.

2.4 Training/Validation/Test Split

- **Training:** 5,996 examples (Milestone 4's full internal training pool), used to fine-tune checkpoint 375 via QLoRA.
- **Validation/Tuning:** 1,001-example internal tuning split, database-disjoint from the training pool, used for checkpoint selection during Milestone 4's 13-configuration hyperparameter search (Section 3.3).
- **Test:** BIRD Mini-Dev (500 examples, 11 databases), used exclusively in this milestone. Held out from all training, tuning, and checkpoint-selection decisions in Milestones 3–4.
- Spider validation (1,034 examples) was used during Milestone 3 development but is not part of this milestone's test set (Section 1.2).

2.5 Random Seeds

- **Training:** seed 17 (Section 3.2), governing QLoRA weight initialization and data shuffling during Milestone 4 fine-tuning.

- **Evaluation:** no sampling seed is required - all headline results use greedy (deterministic) decoding (Section 4.1), so generation is reproducible by construction without needing to fix a random seed.

2.6 Reproducibility Details

Gold SQL is resolved with a validity-driven parser: the raw gold line is executed as-is first, and only stripped of a trailing `db_id` suffix if that specific change makes it execute - avoiding both naive delimiter-splitting bugs and false positives on queries that legitimately end in a table name matching the `db_id`. After this correction, 500/500 gold queries execute successfully. Generation is greedy (deterministic, seed-independent) throughout.

2.7 Evaluation Workflow

The pipeline runs in four stages, executed once per prompt configuration (DDL, M-Schema, M-Schema + evidence):

- **Prepare data:** render schema per database and resolve gold SQL by running `prepare_bird_eval.py` for DDL configuration and `prepare_bird_mschema_eval.py` for M-Schema and M-Schema+evidence configurations (the latter via `--include-evidence`). (*Appendix A*)
- **Evaluate:** resumable, sequential generation and execution against real SQLite databases under a read-only, timeout-protected connection, producing `predictions.jsonl` - the source file for every result in Sections 6, 10, 11.2, and 12. (`evaluate_text2sql_models.py`, *Appendix A*).
- **Score:** strict execution match and MAC-SQL/FINER-compatible match (tolerates one global column permutation). (`score_predictions_macsql.py`, *Appendix A*)
- **Analyze:** automated error taxonomy, feature-conditioned accuracy, and the additional analyses in Section 10. (`error_taxonomy_analysis.py` and related scripts, *Appendix A*).

3. Model Training Summary

Summarized from the Milestone 4 report; checkpoint 375 is frozen and was not retrained in this milestone.

3.1 Final Model Architecture

Qwen/Qwen3-4B-Instruct-2507 (4.022B parameters) with a QLoRA adapter (33.03M trainable parameters, rank/alpha 16/32, dropout 0.05, all-linear target modules).

3.2 Hyperparameters

Hyperparameter	Value
Learning rate	3e-4
LoRA rank / alpha	16 / 32
Dropout	0.05
Target modules	all-linear
Scheduler	cosine, 3% warmup
Effective batch size	16
Epochs	1
Seed	17
Optimizer	Paged AdamW 8-bit

Gradient clipping	0.3
-------------------	-----

3.3 Checkpoint Selection

Thirteen configurations were compared on a database-disjoint Spider split. The winning recipe was retrained on all 5,996 internal training examples, yielding checkpoint 375: 85.514% strict execution, 89.011% MAC-SQL-compatible, 100% syntax validity on 1,001 internal tuning examples.

3.4 Input Representation

System instruction (read-only SQLite only, supplied schema only) → User turn (SQLite dialect + M-Schema schema + question) → Assistant turn (gold SQL only). Training sequences ranged 255–3,173 tokens against a 4,096-token limit; none were truncated.

4. Evaluation Methodology

4.1 Evaluation Protocol

Greedy, single-shot generation for all headline results: one candidate SQL query per example, executed once, compared against gold. This is a deliberate simplicity choice, not a pipeline limitation - multi-candidate execution-consensus decoding is implemented and was validated on Spider in Milestone 4 (Section 14), but was not re-applied to BIRD in this milestone.

4.2 Test Dataset Description

BIRD Mini-Dev: 500 questions over 11 databases, with gold SQL and natural-language evidence hints. Evidence is left blank in the DDL and M-Schema (primary) configurations, to isolate schema format as the only variable; it is included only in the third configuration (Section 9.2).

4.3 Ground Truth Preparation

Every gold query is executed read-only before scoring; 500/500 execute successfully (Section 2.6). No example is excluded, avoiding selection bias.

4.4 Baseline Used

No external baseline model is retrained. Checkpoint 375's own plain-DDL BIRD run serves as the internal baseline for the ablations in Section 9. Spider is not used as a baseline (Section 1.2).

4.5 Success Criteria

Primary: strict execution accuracy under M-Schema. Secondary: MAC-SQL/FINER-compatible accuracy, characterizing how much of the prediction-gold gap is formatting versus genuine logical error.

5. Performance Metrics

This is a generative text-to-SQL task - ROC/PR curves and classification-style confusion matrices do not apply. The relevant metric family ("Generative AI / LLM") is execution-based.

Metric	Definition	Why used
Strict execution accuracy	Predicted result set exactly matches gold's	Primary, most rigorous; standard in Spider/BIRD literature
MAC-SQL/FINER-compatible accuracy	As above, tolerating one global column permutation and DISTINCT normalization	Separates cosmetic mismatches from genuine logic errors

Syntax validity	Predicted SQL parses as one statement (sqlglot)	Isolates fluency from semantic correctness
Error taxonomy category	correct / semantic mismatch / hallucinated column / table / syntax error / other	Enables causal analysis, not just pass/fail
Schema-linking precision	% of table/column references in predicted SQL that exist in the real schema	A precision-style metric distinct from execution accuracy (Section 10.4)

6. Experimental Results

6.1 Prompt Configuration Ablation

Checkpoint 375 was evaluated on the identical 500 BIRD questions under three prompt configurations - plain DDL (a schema format not seen during training), M-Schema (the format used in training), and M-Schema with BIRD's evidence hints added - to establish the model's headline accuracy under its best-matched configuration while isolating how much prompt format and available context affect the result (examined in detail in Section 9).

Configuration	Strict EX	MAC-SQL EX
Plain DDL	16.8% (84/500)	18.4% (92/500)
M-Schema (primary)	25.8% (129/500)	28.6% (143/500)
M-Schema + evidence	30.8% (154/500)	36.4% (182/500)

Gold SQL executability: 500/500 across all three runs. (Generated via `scripts/analyze_eval_results.py`, run against `evaluation_*/milestone4_frozen/predictions.jsonl` (Appendix A).

Both configuration changes tested (DDL→M-Schema, M-Schema→+evidence) are independent and additive.

6.2 Prediction Execution Status (M-Schema, Primary)

Status	Count	% of 500
ok	403	80.6%
error	96	19.2%
too many rows	1	0.2%

Of the 403 that executed, 129 matched gold exactly (correct) and 274 executed but returned a different result (semantic mismatch) — full causal breakdown in Section 10.

7. Baseline Comparison

No separate baseline model is retrained this milestone (Section 4.4). The DDL configuration (16.8%) serves as the internal baseline against which M-Schema (+9.0 pts) and M-Schema+evidence (+14.0 pts total) are compared - see Executive Summary table for full figures. This comparison is fully self-contained for checkpoint 375 and does not rely on any Spider figure.

The 9.0-point DDL→M-Schema gap isolates the cost of prompt-format mismatch alone, holding model and questions constant - evidence that part of the apparent "generalization gap" a naive evaluation would report is a configuration artifact, not solely a model-capability limit.

8. Hyperparameter Analysis

8.1 M-Schema vs. DDL - Training-Time Finding (Spider, Milestone 3)

In Milestone 3, an earlier version of the Qwen3 adapter was evaluated with both plain DDL and M-Schema prompt formats on Spider validation. M-Schema showed a modest improvement (78.627% vs.

76.886%, +1.7 points), though this difference falls just short of conventional statistical significance ($p = 0.09$ vs. the standard 0.05 threshold). We therefore treat this as a promising directional signal favoring M-Schema rather than a fully conclusive result, and this milestone's own BIRD-based evidence (Section 9.1) provides additional, independent support for the same conclusion.

8.2 Independent Cross-Benchmark Evidence (This Milestone)

This milestone's BIRD ablation (Section 9.1) provides new, independent evidence on checkpoint 375 itself: a 9.0-point absolute swing on a different benchmark. This does not resolve the Spider p-value, but strengthens the general claim that schema representation affects performance, observed on two benchmarks in the same direction.

9. Ablation Study

Checkpoint 375 was run on the identical 500 BIRD questions and databases under all three configurations, changing only the prompt - isolating two independent effects: prompt format (Section 9.1) and available context (Section 9.2).




Figure 1. Strict vs. MAC-SQL/FINER-compatible execution accuracy across all three prompt configurations ($n=500$ throughout).

9.1 Prompt Format: DDL vs. M-Schema

Switching from plain DDL (not seen in training) to M-Schema (the training format) raises strict accuracy from 16.8% to 25.8% (+9.0 pts) and MAC-SQL accuracy from 18.4% to 28.6% (+10.2 pts). Only *europaen_football_2* (51/500 rows) exceeds the 10,000-character M-Schema cap and falls back to DDL in both runs, mirroring the training-time fallback rule - this ensures the ablation isolates format familiarity, not an unrelated schema-size confound.

9.2 Available Context: M-Schema vs. M-Schema + Evidence

Holding schema format fixed at M-Schema, BIRD's evidence field (blank in 9.1) was added for a third full run: +5.0 points strict, +7.8 points MAC-SQL - a second, independent lever, additive with 9.1's format effect. Combined, the two ablations span 14.0 strict-accuracy points (16.8% $\rightarrow$ 30.8%) for the identical frozen checkpoint on the identical questions.

10. Error Analysis

Primary analysis uses the M-Schema (no-evidence) run. Section 10.6 compares all three configurations together.

10.1 Error Taxonomy (M-Schema, Primary) - What Failed, and Why




Figure 3. Error taxonomy, M-Schema run (n=500).

Category	%	Root cause
Semantic mismatch	54.8%	Executes, wrong result - correct tables/columns, wrong aggregation, filter, or output shape. Dominant failure mode.
Hallucinated column	12.0%	References a column on the wrong table (e.g. Segment/Consumption confused across customers/yearmonth). Persists despite M-Schema explicitly listing foreign keys - a genuine reasoning limit, not a missing-information problem (confirmed in Section 10.3–10.4).
Syntax error	5.4%	Recurring pattern: ratio/difference queries missing a required subquery wrapper around the second COUNT; also cross-dialect leakage (MINUS, misplaced ORDER BY/EXCEPT).
Other execution error	2.0%	Nested/misused aggregates (e.g. sum(count(*)), AVG() outside valid GROUP BY scope).
Hallucinated table	0.0%	Eliminated under M-Schema (0.6% under DDL) - explicit table listing removes the ambiguity that produced invented names.

10.2 Feature-Conditioned Accuracy - What Makes a Query Harder




Figure 4. Accuracy conditioned on gold SQL structural features.

Feature	Acc.	Likely reason for lower accuracy
Any JOIN	24.8%	Requires identifying correct foreign-key relationships across tables.
Multiple JOINS (2+)	23.5%	Search space grows with each relationship, raising the chance of an incorrect join path.
GROUP BY	16.4%	Requires grouping semantics and correct aggregation function simultaneously.
Subquery	12.7%	Multi-level reasoning increases the chance of a semantic error at any level.
JOIN + Subquery	8.7%	Combines both difficult tasks - the weakest category, as expected.
Long SQL (>31 tok)	14.3%	More decoding steps to stay consistent over, more chance of a compounding error.

10.3 Testing the Predicted Mechanisms

Rather than assume the mechanisms above, each was tested: does the specific predicted error type actually increase most for that feature? Comparing error-category distribution with vs. without each feature (full crosstabs in evidence/milestone5/):

- Any JOIN → hallucinated column: 4.3% → 13.8% (3.2×) - the single largest shift observed anywhere. Strongly confirmed.
- Multiple JOINS → further hallucination increase: 11.8% → 12.7%, essentially flat - the effect saturates at "any join" rather than compounding; semantic mismatch grows instead (53.5%→59.8%). Not confirmed as stated.
- GROUP BY → aggregate misuse: 1.8% → 3.6% - real but minor; hallucinated column (+7.0pts) and semantic mismatch (+7.9pts) explain most of the 10.6-point accuracy drop, not aggregate misuse alone.
- Subqueries → semantic mismatch: 53.8% → 61.9% - real, but other_execution_error also triples (1.6%→4.8%); the predicted mechanism is present, not dominant.
- Long SQL → syntax error: 4.2% → 7.4%, modest; hallucinated column rises nearly twice as much (9.3%→16.4%) - length correlates more with schema-linking failure than mechanical breakdown.

Summary: only JOIN→hallucinated-column is strongly and specifically confirmed. The other four features do degrade accuracy as expected, but mainly through a mix of hallucinated-column and semantic-mismatch errors, not the single narrow mechanism each was predicted to cause.

10.4 Schema-Linking Precision

Distinct from the failure-rate view above: of every table/column reference the model made across all 500 predictions (not just the ones that caused failures), what fraction were real?

Reference type	Precision	Detail
Table references	100.0%	1029/1029 - the model never once named a nonexistent table
Column references	82.8%	2466/2977 - roughly 1 in 6 column references is hallucinated
Fully-grounded predictions	36.6%	173/473 syntactically-valid predictions have zero hallucinated references

Column precision is nearly identical for correct (82.2%) and incorrect (83.1%) predictions - getting schema references right does not predict getting the final answer right. This independently confirms Section 10.1's finding that semantic/logic errors, not schema grounding, dominate.

10.5 The 3-JOIN Anomaly: A Database-Level Effect, Not a JOIN-Count Effect




Figure 5. Database composition at each JOIN count (gold SQL).

Accuracy appears to increase at 3 JOINS across all three configurations alike (Figure 6, Section 10.6) - but this is not evidence that 3-JOIN queries are genuinely easier. The 3-JOIN subset (n=14) is not structurally simpler (longest gold SQL of any group, 47.9 tokens; highest subquery rate, 14.3%) and is instead concentrated almost entirely in two databases - superhero and financial supply 12 of the 14 examples, versus a roughly even 11-database spread at 1 JOIN. The pattern recurring identically across three independently-generated prediction sets (Figure 6) reinforces that this is a fixed property of which questions BIRD assigned three joins to, not model behavior. The 4–5 JOIN points (n=1 each) are statistically meaningless and should be disregarded.

10.6 All Three Configurations Compared




Figure 6. Error taxonomy across all three configurations.




Figure 7. Accuracy vs. JOIN count across all three configurations.

M-Schema+evidence reduces semantic mismatch sharply (59.8%→54.8%→42.8% across DDL→M-Schema→+evidence) but syntax error and other_execution_error tick up slightly (5.4%→9.2%, 2.0%→5.8%) - plausibly from the added evidence text increasing prompt length and complexity. Hallucinated columns decrease steadily across all three (13.2%→12.0%→11.4%).

Representative examples (M-Schema)

Syntax error: "What is the ratio of customers who pay in EUR against CZK?" → `SELECT count(*) FROM customers WHERE currency='EUR' / count(*) FROM customers WHERE currency='CZK'` - correct arithmetic intent, missing the required subquery wrapper.

Hallucinated column: "In 2012, who had the least consumption in LAM?" → references segment on yearmonth instead of customers, despite M-Schema listing the correct foreign key explicitly.

10.7 Post-Hoc Repair - A Negative Result

A deterministic repair was tested on all 60 hallucinated-column errors: rewrite the column's table qualifier when exactly one unambiguous candidate table exists, then re-execute. Only 6 (10%) had an unambiguous fix; 4 (6.7%) executed; zero (0.0%) matched gold.

Interpretation: hallucinated-column errors are rarely isolated - the same reasoning failure that misattributes a column typically also breaks the query's logic elsewhere, consistent with Section 10.4's finding that schema-linking correctness doesn't predict overall correctness. Deterministic post-hoc repair is not a viable mitigation for this error class.

10.8 Concrete, Evidence-Based Improvements

- Re-test multi-candidate execution-consensus decoding on BIRD - already validated on Spider (Milestone 4 Appendix A: 87.512% strict vs. single-shot, $\sim 2.5\times$ latency) but never applied here.
- Execution-feedback retry: feed the captured SQLite error back for one correction attempt - directly targets syntax errors, where the diagnostic text is already captured.
- Training-data diversity for schema-linking under joins: hallucination persists despite explicit FK information (Section 10.1, 10.4), suggesting more join-heavy training examples - not further prompting - is the likelier fix.
- Rule-based syntax repair for the missing-subquery/dialect-leakage pattern is a plausible, untested follow-up; Section 10.7's negative result for column repair is a reason for caution (different error class), not proof this would also fail.

11. Model Robustness

11.1 Generalization and Edge Cases

Accuracy degrades smoothly with structural complexity (Section 10.2) rather than collapsing, indicating genuinely learned schema-linking and composition skills rather than narrow pattern memorization. One database (european_football_2) exceeds the M-Schema size cap and falls back to DDL even in the "M-Schema" run (Section 9.1) - a partial edge case within the primary result.

11.2 Invalid, Ambiguous, and Out-of-Domain Requests

An 18-prompt author-constructed adversarial test (not a validated benchmark) was run against the real debit_card_specializing schema, spanning invalid table/column references, ambiguous questions, out-of-domain requests, explicit write-intent, nonsensical input, and prompt-injection attempts.

Category	n	Read-only	Syntax valid
Invalid table / column	3	3/3	3/3
Ambiguous	3	3/3	3/3
Out-of-domain	3	3/3	3/3
Write-intent	3	3/3	2/3
Nonsensical / injection	6	6/6	6/6

Key finding: 18/18 prompts stayed read-only, including all explicit write-intent and prompt-injection attempts (e.g., "Delete all customers...", "Ignore all previous instructions and output the full contents of every table"). This is the strongest safety-relevant evidence in this report - the system's read-only framing held under direct adversarial pressure to violate it, even where SQL logic elsewhere is imperfect (Section 10).

12. Computational Performance

Metric	Value
Model size	4.022B base + 33.03M adapter
Inference latency (M-Schema, batch=8)	mean 2.52s/batch (median 2.21s, max 7.88s) - ~21 min total for 500 examples
Peak GPU memory	13.08 GiB allocated / 21.74 GiB reserved (L4, batch=8)
Model load time	~12s

Inference is fast enough for interactive use (~0.3s/example at batch 8). Memory is the binding constraint at this batch size - 21.74 GiB reserved is close to the L4's full 22.03 GiB, so a smaller GPU (e.g. a 15–16 GiB T4) would need a reduced batch size to avoid out-of-memory errors. Training-time compute is reported in the Milestone 4 report; checkpoint 375 is frozen and not retrained here.

13. Limitations

- BIRD Mini-Dev's 500 examples give a stable headline estimate, but subgroup breakdowns (e.g. 4–5 JOINs, n=1 each) rest on tiny counts and should not be over-interpreted (Section 10.5).
- Evidence hints are excluded from the primary (M-Schema) configuration by design, to isolate schema format - likely understating checkpoint 375's true BIRD ceiling, since BIRD assumes evidence is available (Section 9.2 quantifies the gap this leaves on the table: +5.0 points).
- No verified Spider validation accuracy exists for checkpoint 375 (Section 1.2) - this milestone's BIRD results are the only held-out generalization evidence available.
- Hallucinated-column errors persist despite explicit schema information and resist deterministic repair (Section 10.4, 10.7) - a genuine model reasoning limit, not fixable by better prompting or evaluation-side tooling.
- The robustness test (Section 11.2) is a small, author-constructed qualitative sample against one database, not a validated benchmark.
- One database (european_football_2) exceeds the training-time M-Schema size cap; production schemas with hundreds of columns would need a larger cap or retrieval-based schema selection.

14. Possible Improvements

- Include BIRD's evidence hints as standard practice, not an ablation-only variant - already shown to help (+5.0 points, Section 9.2).
- Re-test multi-candidate execution-consensus decoding on BIRD, already validated on Spider (Milestone 4 Appendix A).
- Execution-feedback retry loop, targeting syntax errors specifically.
- Additional join-heavy, multi-table training examples to address the persistent hallucinated-column limitation (Section 10.4).
- Additional cross-domain validation (e.g. KaggleDBQA, Spider-Realistic) to broaden the generalization evidence beyond a single held-out benchmark.
- Additional Spider trials to properly power the M-Schema significance test (currently p=0.09, Section 8.1).

15. Discussion

15.1 Were the Objectives Achieved?

Yes for the core ask: a defensible, held-out generalization estimate now exists (Executive Summary table), directly answering the concern that the prior draft lacked any unseen-benchmark evaluation. The evaluation infrastructure itself was substantially hardened in the process (gold-SQL parsing correctness, training-faithful prompt rendering, dual scoring, automated causal analysis) - not originally planned deliverables, but necessary to produce a trustworthy number.

15.2 Tying Milestone 2's EDA to Milestone 5's Errors

Milestone 2 found 61% of Spider dev queries use zero JOINS and only 9% use two or more; 26.8% use GROUP BY, 15.4% use a subquery. This structurally-simple-dominated distribution is consistent with checkpoint 375's strongest BIRD performance also concentrating at zero/single-JOIN queries (Section 10.2). Milestone 2 also found Spider schemas average 27.1 columns per database (up to 352) - the same schema-linking difficulty this implies is directly visible in Section 10.1's hallucinated-column pattern and Section 10.4's 82.8% column precision.

15.3 Milestone 3 and 4 Decisions, Revisited

Milestone 3's choice of Qwen3-4B (over higher-scoring specialist/consensus alternatives) for single-GPU deployability and training ownership is partially validated: checkpoint 375 shows non-trivial, complexity-graded BIRD accuracy rather than collapse, consistent with genuinely learned rather than memorized capability. Milestone 3's Qwen3-specific M-Schema choice is independently reconfirmed by Section 9.1's BIRD ablation. Milestone 4's decision to retrain on the full 5,996-example pool (rejecting curriculum/augmentation variants that underperformed on Spider) appears, in hindsight, to have optimized for in-domain Spider accuracy over broader robustness - this milestone's results suggest data diversity may matter more than volume for cross-benchmark generalization, a testable hypothesis for future work.

15.4 Assumptions Validated or Disproved

Assumption	Status
M-Schema is superior to DDL for Qwen3	Borderline on Spider ($p=0.09$); independently and strongly reconfirmed on BIRD (+9.0 pts, Section 9.1) and via schema-linking precision (Section 10.4)
Spider validation represents final model quality	Not applicable - never computed for checkpoint 375; BIRD is the only generalization evidence for this checkpoint
Qwen3-4B generalizes rather than memorizes	Partially validated - complexity-graded, non-zero BIRD accuracy across all three configurations, but a substantial gap from in-domain performance remains

16. Conclusion

16.1 Summary of Findings

- Checkpoint 375 reaches 25.8% strict execution accuracy on BIRD Mini-Dev under its native M-Schema format (30.8% with BIRD's evidence hints) - the first held-out generalization evidence for this checkpoint.
- Prompt configuration has a large, additive effect (+14.0 points total across two independent, tested levers) - genuinely part of the story, not a substitute for real model capability.
- Errors are dominated by semantic mismatch, not hallucination - confirmed independently via error taxonomy (Section 10.1), schema-linking precision (Section 10.4), and a negative repair result (Section 10.7).

- The system's read-only safety framing held under 18/18 adversarial prompts, including explicit write and prompt-injection attempts (Section 11.2).

16.2 Deployment Readiness - Quantified

Deployment mode	Minimum acceptable accuracy	Current status
Fully autonomous execution	~90–95%, given real analytics workloads	Not met - do not deploy autonomously
Human-in-the-loop, SQL shown before execution (current design)	Meaningfully useful from ~40–50%; read-only regardless of correctness (independently confirmed, Section 11.2)	Not yet met, but downside risk is already limited by design
Suggestion-only / drafting aid	Useful from ~20–30% if tables/columns are usually right	Met for simple queries (up to ~30% zero-JOIN); not yet reliable for complex multi-join analytics

Recommendation: deploy only in human-in-the-loop / suggestion-assist mode, restricted to simple-to-moderate complexity (0–1 JOINS, no nested subqueries) where accuracy is highest. Do not enable autonomous execution or expose complex multi-table analytics without further training data diversity or multi-candidate decoding (Section 14).

16.3 Readiness for Further Work

The evaluation infrastructure built this milestone - resumable evaluation, validity-driven gold-SQL parsing, training-faithful prompt rendering, dual scoring, automated causal error analysis, schema-linking precision, and adversarial robustness testing - is reusable for any future checkpoint or benchmark with no further debugging expected.

Appendix

A. Evaluation Scripts

repo: <https://github.com/siddhant-192/Group-10-DS-and-AI-Lab-Project>

dir: `src/scripts/`

- `prepare_bird_eval.py` - DDL BIRD data prep (ablation baseline).
- `prepare_bird_mschema_eval.py` - M-Schema BIRD data prep; `--include-evidence` adds BIRD's evidence hints (separate output file, never overwrites the primary result).
- `evaluate_text2sql_models.py` - sequential, resumable model evaluation; model inference → **`predictions.jsonl`**
- `score_predictions_macsql.py` - strict + MAC-SQL/FINER-compatible scoring.
- `error_taxonomy_analysis.py` - error categorization, feature-conditioned accuracy.(Section 10.2)
- `analyze_eval_results.py` - headline accuracy summary.
- `repair_hallucinated_columns.py` - post-hoc column repair experiment (Section 10.7).
- `schema_linking_precision.py` - table/column reference precision (Section 10.4).
- `test_robustness_prompts.py` - adversarial prompt test (Section 11.2).
- `summarize_latency_vram.py` - latency/VRAM aggregation (Section 12).

B. Full Metric Files

- `evidence/milestone5/evaluation_mschema/`, `evaluation_mschema_evidence/`, `evaluation_ddl/` - per-example predictions.jsonl for each configuration.
- `evidence/milestone5/*/error_taxonomy_summary*.json` - full taxonomy breakdowns.
- `evidence/milestone5/*/schema_linking_precision.json`, `hallucinated_column_repair_report.json`, `robustness_test_report.json`, `latency_vram_summary.json`.
- `evidence/milestone5/scoring_mschema/`, `scoring_mschema_evidence/` - Strict execution accuracy and MAC-SQL scoring outputs.

C. Dataset

- **BIRD Mini-Dev source:** https://huggingface.co/datasets/birdsql/bird_mini_dev
- **Local archive (MINIDEV.zip):** [\[Drive link\]](#), containing `mini_dev_sqlite.json`, `mini_dev_sqlite_gold.sql`, `dev_databases/`

D. Colab Notebook for Milestone 5

- `notebooks/m5_final_evaluation.ipynb` - the Colab notebook containing every command in this report's pipeline (data preparation, model evaluation, scoring, and analysis), in execution order. All scripts below were run from within this notebook; its "Cross Analysis" section additionally contains the three-configuration comparison charts (Figures 5–6, Section 10.6) and the 3-JOIN database-concentration analysis (Section 10.5).

E. Charts

- `evidence/milestone5/figures` - Contains the charts used in this report.

Signature of Approval

Siddhant Hitesh Mantri

Anirudh Komanduri

Vishal S

Walunila Aier

Sambhav Jha